

Time Series in R: Forecasting and Visualisation

Some automatic forecasting
algorithms

29 May 2017

Outline

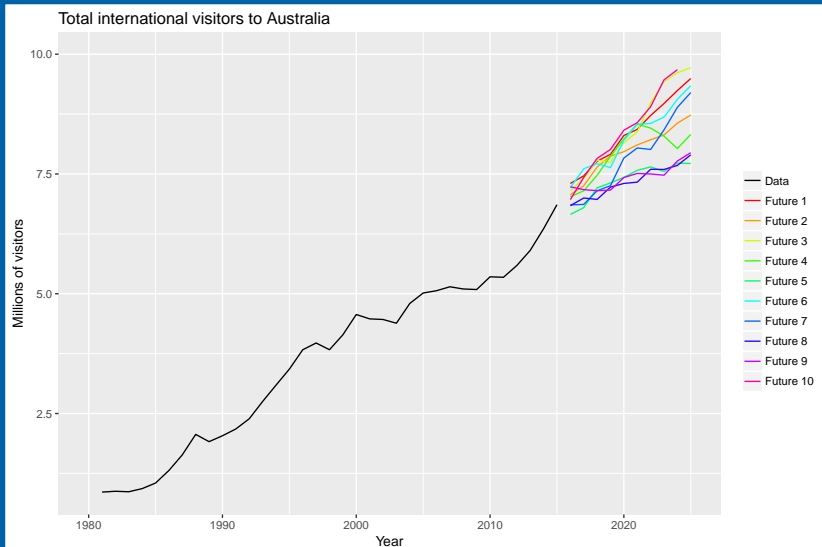
- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 Box-Cox transformations
- 5 ARIMA
- 6 Lab session 6
- 7 STLF

Forecasting


Forecasting is estimating how the sequence of observations will continue into the future.

- We usually think probabilistically about future sample paths
- What range of values covers the possible sample paths with 80% probability?

Sample futures



Forecast intervals



Automatic forecasting algorithms

- 1 Common in business to have over 1000 products that need forecasting at least monthly.
- 2 Forecasts are often required by people who are untrained in time series analysis.

Automatic forecasting algorithms

- 1 Common in business to have over 1000 products that need forecasting at least monthly.
- 2 Forecasts are often required by people who are untrained in time series analysis.


Specifications

Automatic forecasting algorithms must:

- determine an appropriate time series model;
- estimate the parameters;
- compute the forecasts with prediction intervals.

Example: Asian sheep


```
autoplot(livestock)
```



Example: Asian sheep

```
livestock %>% ets %>% forecast %>% autoplot
```


Forecasts from ETS(M,A,N)



Example: Asian sheep


```
livestock %>% auto.arima %>% forecast %>% autoplot
```

Forecasts from ARIMA(0,1,0) with drift



Example: Cortecosteroid sales


```
autoplot(h02)
```



Example: Cortecosteroid sales

```
h02 %>% ets %>% forecast %>% autoplot
```


Forecasts from ETS(M,Ad,M)



Example: Cortecosteroid sales

```
h02 %>% auto.arima %>% forecast %>% autoplot
```


Forecasts from ARIMA(3,1,3)(0,1,1)[12]



Example: Cortecosteroid sales

```
h02 %>% stlf %>% autoplot
```

Forecasts from STL + ETS(M,Ad,N)



Outline

- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 Box-Cox transformations
- 5 ARIMA
- 6 Lab session 6
- 7 STLF

Exponential smoothing methods

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M

N,N: Simple exponential smoothing

A,N: Holt's linear method

A_d,N: Additive damped trend method

A,A: Additive Holt-Winters' method

A,M: Multiplicative Holt-Winters' method

A_d,M: Damped multiplicative Holt-Winters' method

There are also multiplicative trend methods (not recommended).

Exponential smoothing methods

		Seasonal Component		
		N	A	M
Trend Component		(None)	(Additive)	(Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M

- There are 9 separate exponential smoothing methods.
- Each can have an additive or multiplicative error, giving 18 separate models.
- Models with additive and multiplicative errors give the same point forecasts but different prediction intervals.

Exponential smoothing methods

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M

General notation

ETS : ExponenTial Smoothing



Error Trend Seasonal

Exponential smoothing methods

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M

General notation

ETS : Exponential Smoothing



Error Trend Seasonal

ETS(Error,Trend,Seasonal):

- Error = {A,M}
- Trend = {N,A,A_d}
- Seasonal = {N,A,M}.

Automatic ETS forecasting


From Hyndman et al. (IJF, 2002):

- Apply each model that is appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).
- Select best method using AICc:
- Produce forecasts using best method.
- Obtain forecast intervals using underlying state space model.

Method performed very well in M3 competition.

Example: Asian sheep

```
livestock %>% autoplot
```



Example: Asian sheep

```
livestock %>% ets
```

```
## ETS(M,A,N)
##
## Call:
## ets(y = .)
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 1e-04
##
## Initial states:
##   l = 225.1784
##   b = 4.8307
##
## sigma: 0.0344
##
##   AIC  AICc  BIC
## 418.7 420.2 427.9
```

Example: Asian sheep


```
livestock %>% ets %>% forecast
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 2008	460.6	440.3	480.9	429.6	491.6
## 2009	465.4	436.6	494.2	421.3	509.5
## 2010	470.2	434.7	505.8	415.9	524.6
## 2011	475.1	433.8	516.3	412.0	538.1
## 2012	479.9	433.5	526.3	409.0	550.8
## 2013	484.7	433.7	535.8	406.6	562.8
## 2014	489.6	434.1	545.0	404.7	574.4
## 2015	494.4	434.8	554.0	403.2	585.6
## 2016	499.2	435.6	562.8	402.0	596.5
## 2017	504.1	436.7	571.4	401.0	607.1

Example: Asian sheep


```
livestock %>% ets %>% forecast %>% autoplot
```

Forecasts from ETS(M,A,N)



Example: Australian eating-out expenditure

```
ausSAFE %>% autoplot
```



Example: Australian eating-out expenditure

```
auscafe %>% ets
```

```
## ETS(M,A,M)
##
## Call:
## ets(y = .)
##
## Smoothing parameters:
##   alpha = 0.5793
##   beta  = 0.0061
##   gamma = 0.2098
##
## Initial states:
##   l = 0.3458
##   b = 0.0038
##   s=0.9875 0.9452 1.021 1.181 1.026 1.008
##         0.9728 0.9793 0.9796 0.9379 0.9931 0.9686
##
## sigma: 0.0245
##
##      AIC      AICc      BIC
## -339.1 -337.6 -270.6
```

Example: Australian eating-out expenditure


```
auscfe %>% ets %>% forecast
```

##		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	Dec 2016	4.035	3.908	4.162	3.841	4.229
##	Jan 2017	3.637	3.505	3.769	3.435	3.839
##	Feb 2017	3.331	3.195	3.467	3.123	3.539
##	Mar 2017	3.693	3.528	3.859	3.440	3.946
##	Apr 2017	3.656	3.478	3.833	3.384	3.928
##	May 2017	3.701	3.508	3.894	3.406	3.996
##	Jun 2017	3.591	3.392	3.790	3.286	3.896
##	Jul 2017	3.823	3.599	4.047	3.480	4.166
##	Aug 2017	3.841	3.604	4.079	3.479	4.204
##	Sep 2017	3.798	3.552	4.044	3.422	4.174
##	Oct 2017	3.856	3.596	4.117	3.458	4.255
##	Nov 2017	3.827	3.558	4.096	3.415	4.239
##	Dec 2017	4.194	3.878	4.509	3.711	4.676
##	Jan 2018	3.780	3.486	4.074	3.330	4.229
##	Feb 2018	3.461	3.183	3.739	3.036	3.886
##	Mar 2018	3.837	3.520	4.155	3.351	4.323

Example: Australian eating-out expenditure


```
ausafe %>% ets %>% forecast %>% autoplot
```

Forecasts from ETS(M,A,M)



Example: Half-hourly electricity demand

```
taylor %>% autoplot
```



Example: Half-hourly electricity demand

```
taylor %>% ets
```

```
## Warning in ets(.): I can't handle data with
## frequency greater than 24. Seasonality will
## be ignored. Try stlf() if you need seasonal
## forecasts.
```

```
## ETS(A,Ad,N)
```

```
##
```

```
## Call:
```

```
##   ets(y = .)
```

```
##
```

```
##   Smoothing parameters:
```

```
##     alpha = 0.9999
```

```
##     beta  = 0.9999
```

```
##     phi   = 0.8696
```

```
##
```

```
##   Initial states:
```

```
##     l = 22509.4085
```

```
##     s = 250.5000
```

Outline

- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 Box-Cox transformations
- 5 ARIMA
- 6 Lab session 6
- 7 STLF

Lab Session 5

Outline

- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 **Box-Cox transformations**
- 5 ARIMA
- 6 Lab session 6
- 7 STLF

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as y_1, \dots, y_n and transformed observations as w_1, \dots, w_n .

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as y_1, \dots, y_n and transformed observations as w_1, \dots, w_n .

Mathematical transformations for stabilizing variation

Square root	$w_t = \sqrt{y_t}$	↓
Cube root	$w_t = \sqrt[3]{y_t}$	Increasing
Logarithm	$w_t = \log(y_t)$	strength

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as y_1, \dots, y_n and transformed observations as w_1, \dots, w_n .


Mathematical transformations for stabilizing variation

Square root	$w_t = \sqrt{y_t}$	↓
Cube root	$w_t = \sqrt[3]{y_t}$	Increasing
Logarithm	$w_t = \log(y_t)$	strength

Logarithms, in particular, are useful because they are more interpretable: changes in a log value are **relative (percent) changes on the original scale**.


Variance stabilization

Square root electricity production




Variance stabilization

Cube root electricity production




Variance stabilization

Log electricity production



Variance stabilization

Inverse electricity production



Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:

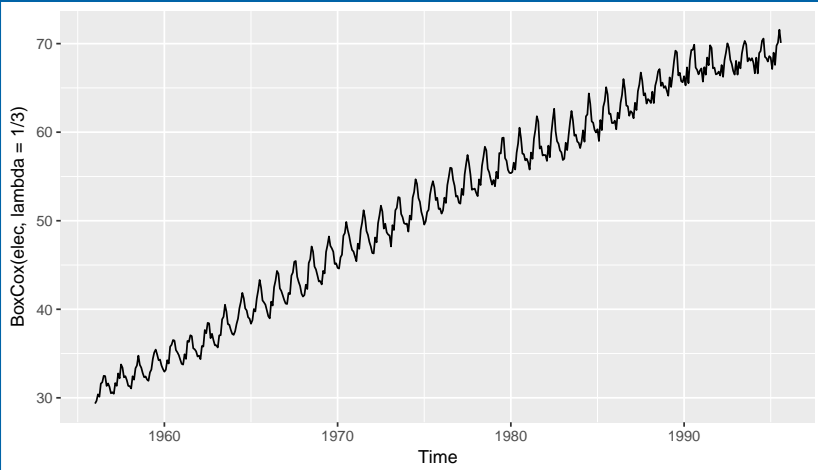
$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- $\lambda = 1$: (No substantive transformation)
- $\lambda = \frac{1}{2}$: (Square root plus linear transformation)
- $\lambda = 0$: (Natural logarithm)
- $\lambda = -1$: (Inverse plus 1)

Box-Cox transformations

Box-Cox transformations

```
autoplot(BoxCox(elec,lambda=1/3))
```



Automated Box-Cox transformations

```
(BoxCox.lambda(elec))
```

```
## [1] 0.2654
```

Automated Box-Cox transformations

```
(BoxCox.lambda(elec))
```

```
## [1] 0.2654
```

- This attempts to balance the seasonal fluctuations and random variation across the series.
- Always check the results.
- A low value of λ can give extremely large prediction intervals.

Outline

- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 Box-Cox transformations
- 5 **ARIMA**
- 6 Lab session 6
- 7 STLF

How does auto.arima() work?

Non-seasonal version:

A non-seasonal ARIMA process

$$\phi(B)(1 - B)^d y_t = c + \theta(B)\varepsilon_t$$

Need to select appropriate orders: p, q, d .

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via unit root tests.
- Select p, q by minimising AICc.
- Use stepwise search to traverse model space.

How does `auto.arima()` work?

Non-seasonal version:

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

How does auto.arima() work?

Non-seasonal version:

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:


- vary one of p, q , from current model by ± 1 ;
- p, q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

Repeat Step 2 until no lower AICc can be found.

Example: Asian sheep

```
livestock %>% autoplot
```



Example: Asian sheep


```
livestock %>% auto.arima
```

```
## Series: .  
## ARIMA(0,1,0) with drift  
##  
## Coefficients:  
##      drift  
##      4.858  
## s.e.  1.789  
##  
## sigma^2 estimated as 150:  log likelihood=-180.1  
## AIC=364.2   AICc=364.4   BIC=367.8
```

Example: Asian sheep


```
livestock %>% auto.arima %>% forecast %>% autoplot
```

Forecasts from ARIMA(0,1,0) with drift




Example: Australian eating-out expenditure

```
ausSAFE %>% autoplot
```



Example: Australian eating-out expenditure

```
ausSAFE %>% BoxCox(lambda=0) %>% autoplot
```



Example: Australian eating-out expenditure


```
auscafe %>% auto.arima(lambda=0)
```

```
## Series: .  
## ARIMA(1,1,2)(0,0,2)[12] with drift  
## Box Cox transformation: lambda= 0  
##  
## Coefficients:  
##          ar1      ma1      ma2      sma1      sma2      drift  
##        -0.787   0.516  -0.389   0.746   0.474   0.006  
## s.e.    0.051   0.061   0.046   0.054   0.039   0.002  
##  
## sigma^2 estimated as 0.0013:  log likelihood=788.8  
## AIC=-1564   AICc=-1563   BIC=-1535
```

Example: Australian eating-out expenditure

```
auscfe %>% auto.arima(lambda=0) %>%  
  forecast %>% autoplot
```

Forecasts from ARIMA(1,1,2)(0,0,2)[12] with drift



Example: Australian eating-out expenditure


```
auscfe %>% auto.arima(lambda=0, D=1)
```

```
## Series: .  
## ARIMA(3,1,0)(2,1,1)[12]  
## Box Cox transformation: lambda= 0  
##  
## Coefficients:  
##          ar1      ar2      ar3      sar1      sar2  
##        -0.340  -0.108   0.095   0.125  -0.059  
## s.e.      0.052   0.053   0.050   0.066   0.059  
##          sma1  
##        -0.828  
## s.e.      0.044  
##  
## sigma^2 estimated as 0.000572:  log likelihood=929.5  
## AIC=-1845   AICc=-1845   BIC=-1817
```

Example: Australian eating-out expenditure


```
auscafe %>% auto.arima(lambda=0, D=1) %>%  
  forecast %>% autoplot
```

Forecasts from ARIMA(3,1,0)(2,1,1)[12]



Example: Half-hourly electricity demand

```
taylor %>% autoplot
```



Example: Half-hourly electricity demand


```
taylor %>% auto.arima
```

```
## Series: .  
## ARIMA(4,0,5) with non-zero mean  
##  
## Coefficients:  
##          ar1      ar2      ar3      ar4      ma1      ma2  
##          1.726  -0.378  -0.553   0.190   0.391  -0.075  
## s.e.    0.080   0.180   0.151   0.053   0.080   0.052  
##          ma3      ma4      ma5      mean  
##          -0.332  -0.255  -0.177  29611.3  
## s.e.    0.019   0.032   0.021   230.6  
##  
## sigma^2 estimated as 159426:  log likelihood=-29870  
## AIC=59762   AICc=59762   BIC=59832
```

Example: Half-hourly electricity demand

```
taylor %>% auto.arima %>% forecast %>% autoplot
```

Forecasts from ARIMA(4,0,5) with non-zero mean



Outline

- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 Box-Cox transformations
- 5 ARIMA
- 6 Lab session 6
- 7 STLF

Lab Session 6

Outline


- 1 Automatic forecasting
- 2 ETS
- 3 Lab session 5
- 4 Box-Cox transformations
- 5 ARIMA
- 6 Lab session 6
- 7 **STLF**

STLF

- Decomposes time series into a trend, seasonal and remainder component using STL decomposition
- Forecasts the seasonally adjusted series using ETS
- Forecasts the seasonal component using a “seasonal naive” approach (replicating last available year).
- Combines them to get forecasts for original series.
- May need a Box-Cox transformation


Example: Australian eating-out expenditure

```
ausSAFE %>% autoplot
```




Example: Australian eating-out expenditure

```
auscape %>% BoxCox(lambda=0) %>% autoplot
```



Example: Australian eating-out expenditure


```
auscafe %>% BoxCox(lambda=0) %>%  
  stl(s.window='periodic') %>% autoplot
```



Example: Australian eating-out expenditure

```
auscafe %>% stlf(lambda=0) %>% autoplot
```

Forecasts from STL + ETS(A,A,N)



Example: Half-hourly electricity demand

```
taylor %>% stlf %>% autoplot
```

Forecasts from STL + ETS(M,N,N)

