

Time series and forecasting in R

Rob J Hyndman

29 June 2008



MONASH University

Outline

- 1 Time series objects
- 2 Basic time series functionality
- 3 The forecast package
- 4 Exponential smoothing
- 5 ARIMA modelling
- 6 More from the forecast package
- 7 Time series packages on CRAN

Outline

- 1 **Time series objects**
- 2 Basic time series functionality
- 3 The forecast package
- 4 Exponential smoothing
- 5 ARIMA modelling
- 6 More from the forecast package
- 7 Time series packages on CRAN

Australian GDP

```
ausgdp <- ts(scan("gdp.dat"),frequency=4,  
             start=1971+2/4)
```

Australian GDP

```
ausgdp <- ts(scan("gdp.dat"),frequency=4,  
             start=1971+2/4)
```

- Class: `ts`
- Print and plotting methods available.

Australian GDP


```
ausgdp <- ts(scan("gdp.dat"),frequency=4,  
             start=1971+2/4)
```

- Class: `ts`
- Print and plotting methods available.

```
> ausgdp
```

	Qtr1	Qtr2	Qtr3	Qtr4
1971			4612	4651
1972	4645	4615	4645	4722
1973	4780	4830	4887	4933
1974	4921	4875	4867	4905
1975	4938	4934	4942	4979
1976	5028	5079	5112	5127
1977	5130	5101	5072	5069
1978	5100	5166	5244	5312

Australian GDP



Australian beer production

```
> beer
```


	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1991	164	148	152	144	155	125	153	146	138	190	192	192
1992	147	133	163	150	129	131	145	137	138	168	176	188
1993	139	143	150	154	137	129	128	140	143	151	177	184
1994	151	134	164	126	131	125	127	143	143	160	190	182
1995	138	136	152	127	151	130	119	153				

Outline

- 1 Time series objects
- 2 Basic time series functionality**
- 3 The forecast package
- 4 Exponential smoothing
- 5 ARIMA modelling
- 6 More from the forecast package
- 7 Time series packages on CRAN

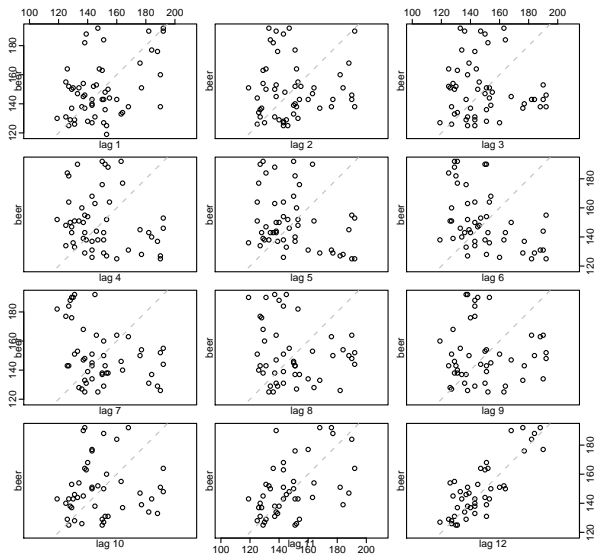
Lag plots

```
> lag.plot(beer, lags=12)
```



Lag plots


```
> lag.plot(beer, lags=12, do.lines=FALSE)
```




Lag plots

```
lag.plot(x, lags = 1, layout = NULL,  
  set.lags = 1:lags, main = NULL,  
  asp = 1, diag = TRUE,  
  diag.col = "gray", type = "p",  
  oma = NULL, ask = NULL,  
  do.lines = (n <= 150), labels = do.lines,  
  ...)
```

ACF



PACF




ACF/PACF

```
acf(x, lag.max = NULL,  
    type = c("correlation", "covariance", "partial"),  
    plot = TRUE, na.action = na.fail, demean = TRUE, ...)
```

```
pacf(x, lag.max, plot, na.action, ...)
```

```
ARMAacf(ar = numeric(0), ma = numeric(0), lag.max = r,  
        pacf = FALSE)
```



Spectrum



Spectrum

AR(12) spectrum

```
> spectrum(beer, method="ar")
```



Spectrum

```
spectrum(x, ..., method = c("pgram", "ar"))
```


```
spec.pgram(x, spans = NULL, kernel, taper = 0.1,  
  pad = 0, fast = TRUE, demean = FALSE,  
  detrend = TRUE, plot = TRUE,  
  na.action = na.fail, ...)
```

```
spec.ar(x, n.freq, order = NULL, plot = TRUE,  
  na.action = na.fail,  
  method = "yule-walker", ...)
```

Classical decomposition


`decompose(beer)`

Decomposition of additive time series



STL decomposition

```
plot(stl(beer,s.window="periodic"))
```



Decomposition

```
decompose(x, type = c("additive", "multiplicative"),  
          filter = NULL)
```

```
stl(x, s.window, s.degree = 0,  
     t.window = NULL, t.degree = 1,  
     l.window = nextodd(period), l.degree = t.degree,  
     s.jump = ceiling(s.window/10),  
     t.jump = ceiling(t.window/10),  
     l.jump = ceiling(l.window/10),  
     robust = FALSE,  
     inner = if(robust) 1 else 2,  
     outer = if(robust) 15 else 0,  
     na.action = na.fail)
```

Outline

- 1 Time series objects
- 2 Basic time series functionality
- 3 The forecast package**
- 4 Exponential smoothing
- 5 ARIMA modelling
- 6 More from the forecast package
- 7 Time series packages on CRAN


forecast package

```
> forecast(beer)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Sep 1995	138.5042	128.2452	148.7632	122.8145	154.1940
Oct 1995	169.1987	156.6506	181.7468	150.0081	188.3894
Nov 1995	181.6725	168.1640	195.1810	161.0131	202.3320
Dec 1995	178.5394	165.2049	191.8738	158.1461	198.9327
Jan 1996	144.0816	133.2492	154.9140	127.5148	160.6483
Feb 1996	135.7967	125.4937	146.0996	120.0396	151.5537
Mar 1996	151.4813	139.8517	163.1110	133.6953	169.2673
Apr 1996	138.9345	128.1106	149.7584	122.3808	155.4882
May 1996	138.5279	127.5448	149.5110	121.7307	155.3250
Jun 1996	127.0269	116.7486	137.3052	111.3076	142.7462
Jul 1996	134.9452	123.7716	146.1187	117.8567	152.0337
Aug 1996	145.3088	132.9658	157.6518	126.4318	164.1858
Sep 1996	139.7348	127.4679	152.0018	120.9741	158.4955
Oct 1996	170.6709	155.2397	186.1020	147.0709	194.2708
Nov 1996	183.2204	166.1298	200.3110	157.0826	209.3582
Dec 1996	180.0290	162.6798	197.3783	153.4957	206.5624
Jan 1997	145.2589	130.7803	159.7374	123.1159	167.4019
Feb 1997	136.8222	122.7525	151.0071	115.8822	159.4822

forecast package

Forecasts from ETS(M,Ad,M)



forecast package

```
> summary(forecast(beer))
```

Forecast method: ETS(M,Ad,M)

Smoothing parameters:

alpha = 0.0267

beta = 0.0232

gamma = 0.025

phi = 0.98

Initial states:

l = 162.5752

b = -0.1598

s = 1.1979 1.2246 1.1452 0.9354 0.9754 0.9068

0.8523 0.9296 0.9342 1.0160 0.9131 0.9696

sigma: 0.0578

AIC

AICc

BIC

400.0005 515.1247 522.4604

forecast package

- Automatic exponential smoothing state space modelling.

forecast package

- Automatic exponential smoothing state space modelling.
- Automatic ARIMA modelling

forecast package

- Automatic exponential smoothing state space modelling.
- Automatic ARIMA modelling
- Forecasting intermittent demand data using Croston's method

forecast package

- Automatic exponential smoothing state space modelling.
- Automatic ARIMA modelling
- Forecasting intermittent demand data using Croston's method
- Forecasting using Theta method

forecast package

- Automatic exponential smoothing state space modelling.
- Automatic ARIMA modelling
- Forecasting intermittent demand data using Croston's method
- Forecasting using Theta method
- Forecasting methods for most time series modelling functions including `arima()`, `ar()`, `StructTS()`, `ets()`, and others.

forecast package


- Automatic exponential smoothing state space modelling.
- Automatic ARIMA modelling
- Forecasting intermittent demand data using Croston's method
- Forecasting using Theta method
- Forecasting methods for most time series modelling functions including `arima()`, `ar()`, `StructTS()`, `ets()`, and others.
- Part of the **forecasting** bundle along with **fma**, **expsmooth** and **Mcomp**.

Outline

- 1 Time series objects
- 2 Basic time series functionality
- 3 The forecast package
- 4 Exponential smoothing**
- 5 ARIMA modelling
- 6 More from the forecast package
- 7 Time series packages on CRAN

Exponential smoothing


Classic Reference



Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, 3rd ed., Wiley: NY.

Exponential smoothing

Classic Reference



Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, 3rd ed., Wiley: NY.

Current Reference



Hyndman, Koehler, Ord and Snyder (2008) *Forecasting with exponential smoothing: the state space approach*, Springer-Verlag: Berlin.

Exponential smoothing

- Until recently, there has been no stochastic modelling framework incorporating likelihood calculation, prediction intervals, etc.

Exponential smoothing

- Until recently, there has been no stochastic modelling framework incorporating likelihood calculation, prediction intervals, etc.
- Ord, Koehler & Snyder (JASA, 1997) and Hyndman, Koehler, Snyder and Grose (IJF, 2002) showed that all ES methods (including non-linear methods) are optimal forecasts from innovation state space models.

Exponential smoothing

- Until recently, there has been no stochastic modelling framework incorporating likelihood calculation, prediction intervals, etc.
- Ord, Koehler & Snyder (JASA, 1997) and Hyndman, Koehler, Snyder and Grose (IJF, 2002) showed that all ES methods (including non-linear methods) are optimal forecasts from innovation state space models.
- Hyndman et al. (2008) provides a comprehensive and up-to-date survey of the area.

Exponential smoothing

- Until recently, there has been no stochastic modelling framework incorporating likelihood calculation, prediction intervals, etc.
- Ord, Koehler & Snyder (JASA, 1997) and Hyndman, Koehler, Snyder and Grose (IJF, 2002) showed that all ES methods (including non-linear methods) are optimal forecasts from innovation state space models.
- Hyndman et al. (2008) provides a comprehensive and up-to-date survey of the area.
- The **forecast** package implements the framework of HKSO.

Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

General notation **ETS**(*Error, Trend, Seasonal*)

Exponential smoothing

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

Exponential smoothing

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

ETS(A,N,N): Simple exponential smoothing with additive errors

Exponential smoothing

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

ETS(A,A,N): Holt's linear method with additive errors

Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

ETS(A,A,A): Additive Holt-Winters' method with additive errors

Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A _d	(Additive damped)	A _d ,N	A _d ,A	A _d ,M
M	(Multiplicative)	M,N	M,A	M,M
M _d	(Multiplicative damped)	M _d ,N	M _d ,A	M _d ,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

ETS(M,A,M): Multiplicative Holt-Winters' method with multiplicative errors

Exponential smoothing

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A_d	(Additive damped)	A_d,N	A_d,A	A_d,M
M	(Multiplicative)	M,N	M,A	M,M
M_d	(Multiplicative damped)	M_d,N	M_d,A	M_d,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

ETS(A, A_d , N): Damped trend method with additive errors

Exponential smoothing

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A_d	(Additive damped)	A_d,N	A_d,A	A_d,M
M	(Multiplicative)	M,N	M,A	M,M
M_d	(Multiplicative damped)	M_d,N	M_d,A	M_d,M

General notation **ETS**(*Error, Trend, Seasonal*)
ExponenTial Smoothing

There are 30 separate models in the ETS framework

Innovations state space models

**No trend or seasonality
and multiplicative errors**

Example: ETS(M,N,N)

$$y_t = l_{t-1}(1 + \varepsilon_t)$$

$$\begin{aligned} l_t &= \alpha y_t + (1 - \alpha)l_{t-1} \\ &= l_{t-1}(1 + \alpha\varepsilon_t) \end{aligned}$$

$$0 \leq \alpha \leq 1$$

ε_t is white noise with mean zero.

Innovations state space models

**No trend or seasonality
and multiplicative errors**

Example: ETS(M,N,N)

$$y_t = \ell_{t-1}(1 + \varepsilon_t)$$

$$\begin{aligned}\ell_t &= \alpha y_t + (1 - \alpha)\ell_{t-1} \\ &= \ell_{t-1}(1 + \alpha\varepsilon_t)\end{aligned}$$

$$0 \leq \alpha \leq 1$$

ε_t is white noise with mean zero.

**All exponential smoothing models can be
written using analogous state space equations.**

Innovation state space models

Let $\mathbf{x}_t = (\ell_t, b_t, s_t, s_{t-1}, \dots, s_{t-m+1})$ and $\varepsilon_t \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$.

Example: Holt-Winters' multiplicative seasonal method

Example: ETS(M,A,M)

$$Y_t = (\ell_{t-1} + b_{t-1})s_{t-m}(1 + \varepsilon_t)$$

$$\ell_t = \alpha(y_t/s_{t-m}) + (1 - \alpha)(\ell_{t-1} + b_{t-1})$$

$$b_t = \beta(\ell_t - \ell_{t-1}) + (1 - \beta)b_{t-1}$$

$$s_t = \gamma(y_t/(\ell_{t-1} + b_{t-1})) + (1 - \gamma)s_{t-m}$$

where $0 \leq \alpha \leq 1$, $0 \leq \beta \leq \alpha$, $0 \leq \gamma \leq 1 - \alpha$
and m is the period of seasonality.

Exponential smoothing

From Hyndman et al. (2008):

- Apply each of 30 methods that are appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).

Exponential smoothing

From Hyndman et al. (2008):

- Apply each of 30 methods that are appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).
- Select best method using AIC:

$$\text{AIC} = -2 \log(\text{Likelihood}) + 2p$$

where $p = \#$ parameters.

Exponential smoothing

From Hyndman et al. (2008):

- Apply each of 30 methods that are appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).
- Select best method using AIC:

$$\text{AIC} = -2 \log(\text{Likelihood}) + 2p$$

where $p = \#$ parameters.

- Produce forecasts using best method.

Exponential smoothing

From Hyndman et al. (2008):

- Apply each of 30 methods that are appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).
- Select best method using AIC:

$$\text{AIC} = -2 \log(\text{Likelihood}) + 2p$$

where $p = \#$ parameters.

- Produce forecasts using best method.
- Obtain prediction intervals using underlying state space model.

Exponential smoothing

From Hyndman et al. (2008):

- Apply each of 30 methods that are appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).
- Select best method using AIC:

$$\text{AIC} = -2 \log(\text{Likelihood}) + 2p$$

where $p = \#$ parameters.

- Produce forecasts using best method.
- Obtain prediction intervals using underlying state space model.

Exponential smoothing

From Hyndman et al. (2008):

- Apply each of 30 methods that are appropriate to the data. Optimize parameters and initial values using MLE (or some other criterion).
- Select best method using AIC:

$$\text{AIC} = -2 \log(\text{Likelihood}) + 2p$$

where $p = \#$ parameters.

- Produce forecasts using best method.
- Obtain prediction intervals using underlying state space model.

Method performed very well in M3 competition.

Exponential smoothing

```
fit <- ets(beer)
fit2 <- ets(beer,model="MNM",damped=FALSE)
fcast1 <- forecast(fit, h=24)
fcast2 <- forecast(fit2, h=24)
```

Exponential smoothing

```
fit <- ets(beer)
fit2 <- ets(beer,model="MNM",damped=FALSE)
fcast1 <- forecast(fit, h=24)
fcast2 <- forecast(fit2, h=24)
```

```
ets(y, model="ZZZ", damped=NULL, alpha=NULL, beta=NULL,
    gamma=NULL, phi=NULL, additive.only=FALSE,
    lower=c(rep(0.01,3), 0.8), upper=c(rep(0.99,3),0.98),
    opt.crit=c("lik","amse","mse","sigma"), nmse=3,
    bounds=c("both","usual","admissible"),
    ic=c("aic","aicc","bic"), restrict=TRUE)
```

Exponential smoothing

```
> fit  
ETS(M,Ad,M)
```

Smoothing parameters:

alpha = 0.0267

beta = 0.0232

gamma = 0.025

phi = 0.98

Initial states:

l = 162.5752

b = -0.1598

s = 1.1979 1.2246 1.1452 0.9354 0.9754 0.9068

0.8523 0.9296 0.9342 1.016 0.9131 0.9696

sigma: 0.0578

AIC	AICc	BIC
499.0295	515.1347	533.4604

Exponential smoothing

```
> fit2  
ETS(M,N,M)
```

```
Smoothing parameters:
```

```
alpha = 0.247
```

```
gamma = 0.01
```

```
Initial states:
```

```
l = 168.1208
```

```
s = 1.2417 1.2148 1.1388 0.9217 0.9667 0.8934
```

```
0.8506 0.9182 0.9262 1.049 0.9047 0.9743
```

```
sigma: 0.0604
```

AIC	AICc	BIC
500.0439	510.2878	528.3988

Exponential smoothing

ets() function

- Automatically chooses a model by default using the AIC

Exponential smoothing

ets() function

- Automatically chooses a model by default using the AIC
- Can handle any combination of trend, seasonality and damping

Exponential smoothing

ets() function

- Automatically chooses a model by default using the AIC
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model

Exponential smoothing

ets() function

- Automatically chooses a model by default using the AIC
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model
- Ensures the parameters are admissible (equivalent to invertible)

Exponential smoothing

ets() function

- Automatically chooses a model by default using the AIC
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model
- Ensures the parameters are admissible (equivalent to invertible)
- Produces an object of class `ets`.

Exponential smoothing

ets objects

- **Methods:** `coef()`, `plot()`,
`summary()`, `residuals()`, `fitted()`,
`simulate()` and `forecast()`


Exponential smoothing

ets objects

- **Methods:** `coef()`, `plot()`, `summary()`, `residuals()`, `fitted()`, `simulate()` and `forecast()`
- `plot()` function shows time plots of the original time series along with the extracted components (level, growth and seasonal).

Exponential smoothing

`plot(fit)`
Decomposition by ETS(M,Ad,M) method



Goodness-of-fit

```
> accuracy(fit)
```


ME	RMSE	MAE	MPE	MAPE	MASE
0.0774	8.4156	7.0331	-0.2915	4.7883	0.4351

```
> accuracy(fit2)
```

ME	RMSE	MAE	MPE	MAPE	MASE
-1.3884	9.0015	7.3303	-1.1945	5.0237	0.4535


Forecast intervals

Forecasts from ETS(M,Ad,M)



Forecast intervals

Forecasts from ETS(M,Ad,M)



Exponential smoothing

`ets()` function also allows refitting model to new data set.

```
> usfit <- ets(usnetelec[1:45])  
> test <- ets(usnetelec[46:55], model = usfit)
```

```
> accuracy(test)
```

ME	RMSE	MAE	MPE	MAPE	MASE
-4.3057	58.1668	43.5241	-0.1023	1.1758	0.5206

```
> accuracy(forecast(usfit,10), usnetelec[46:55])
```

ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
46.36580	65.55163	49.83883	1.25087	1.35781	0.72895	0.08899	0.73725

forecast package

forecast() function

- Takes either a time series as its main argument, or a time series model.

forecast package

forecast() function

- Takes either a time series as its main argument, or a time series model.
- Methods for objects of class `ts`, `ets`, `arima`, `HoltWinters`, `StructTS`, `ar` and others.

forecast package

forecast() function

- Takes either a time series as its main argument, or a time series model.
- Methods for objects of class `ts`, `ets`, `arima`, `HoltWinters`, `StructTS`, `ar` and others.
- If argument is `ts`, it uses `ets` model.

forecast package

forecast() function

- Takes either a time series as its main argument, or a time series model.
- Methods for objects of class `ts`, `ets`, `arima`, `HoltWinters`, `StructTS`, `ar` and others.
- If argument is `ts`, it uses `ets` model.
- Calls `predict()` when appropriate.

forecast package

forecast() function

- Takes either a time series as its main argument, or a time series model.
- Methods for objects of class `ts`, `ets`, `arima`, `HoltWinters`, `StructTS`, `ar` and others.
- If argument is `ts`, it uses `ets` model.
- Calls `predict()` when appropriate.
- Output as class `forecast`.

forecast package

forecast class contains

- Original series
- Point forecasts
- Prediction intervals
- Forecasting method used
- Forecasting model information
- Residuals
- One-step forecasts for observed data

forecast package

forecast class contains

- Original series
- Point forecasts
- Prediction intervals
- Forecasting method used
- Forecasting model information
- Residuals
- One-step forecasts for observed data

Methods applying to the forecast class:

- print
- plot
- summary

Outline

- 1 Time series objects
- 2 Basic time series functionality
- 3 The forecast package
- 4 Exponential smoothing
- 5 ARIMA modelling**
- 6 More from the forecast package
- 7 Time series packages on CRAN

ARIMA modelling

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.

ARIMA modelling

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary

ARIMA modelling

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`

ARIMA modelling

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`
- It does not allow re-fitting a model to new data.

ARIMA modelling

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`
- It does not allow re-fitting a model to new data.
- So I prefer the `Arima()` function in the **forecast** package which acts as a wrapper to `arima()`.

ARIMA modelling

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`
- It does not allow re-fitting a model to new data.
- So I prefer the `Arima()` function in the **forecast** package which acts as a wrapper to `arima()`.
- Even better, the `auto.arima()` function in the **forecast** package.

ARIMA modelling

```
> fit <- auto.arima(beer)
> fit
Series: beer
ARIMA(0,0,0)(1,0,0)[12] with non-zero mean
```

Coefficients:

	sar1	intercept
	0.8431	152.1132
s.e.	0.0590	5.1921

sigma² estimated as 122.1: log likelihood = -221.44
AIC = 448.88 AICc = 449.34 BIC = 454.95

How does `auto.arima()` work?

A seasonal ARIMA process

$$\Phi(B^m)\phi(B)(1 - B^m)^D(1 - B)^d y_t = c + \Theta(B^m)\theta(B)\varepsilon_t$$

Need to select appropriate orders: p, q, P, Q, D, d

How does `auto.arima()` work?

A seasonal ARIMA process

$$\Phi(B^m)\phi(B)(1 - B^m)^D(1 - B)^d y_t = c + \Theta(B^m)\theta(B)\varepsilon_t$$

Need to select appropriate orders: p, q, P, Q, D, d

Use Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d and D via unit root tests.
- Select p, q, P, Q by minimising AIC.
- Use stepwise search to traverse model space.

How does `auto.arima()` work?

$$\text{AIC} = -2\log(L) + 2(p + q + P + Q + k)$$

where L is the maximised likelihood fitted to the *differenced* data,
 $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

How does `auto.arima()` work?

$$\text{AIC} = -2\log(L) + 2(p + q + P + Q + k)$$

where L is the maximised likelihood fitted to the *differenced* data,
 $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Step 1: Select current model (with smallest AIC) from:

ARIMA(2, d , 2)(1, D , 1) _{m}

ARIMA(0, d , 0)(0, D , 0) _{m}

ARIMA(1, d , 0)(1, D , 0) _{m}

ARIMA(0, d , 1)(0, D , 1) _{m}

if seasonal

How does `auto.arima()` work?

$$AIC = -2\log(L) + 2(p + q + P + Q + k)$$

where L is the maximised likelihood fitted to the *differenced* data,
 $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Step 1: Select current model (with smallest AIC) from:

ARIMA(2, d , 2)(1, D , 1) _{m}

ARIMA(0, d , 0)(0, D , 0) _{m}

ARIMA(1, d , 0)(1, D , 0) _{m} if seasonal

ARIMA(0, d , 1)(0, D , 1) _{m}

Step 2: Consider variations of current model:

- vary one of p, q, P, Q from current model by ± 1
- p, q both vary from current model by ± 1 .
- P, Q both vary from current model by ± 1 .
- Include/exclude c from current model

Model with lowest AIC becomes current model.

How does `auto.arima()` work?

$$\text{AIC} = -2\log(L) + 2(p + q + P + Q + k)$$

where L is the maximised likelihood fitted to the *differenced* data,
 $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Step 1: Select current model (with smallest AIC) from:

ARIMA(2, d , 2)(1, D , 1) _{m}

ARIMA(0, d , 0)(0, D , 0) _{m}

ARIMA(1, d , 0)(1, D , 0) _{m} if seasonal

ARIMA(0, d , 1)(0, D , 1) _{m}

Step 2: Consider variations of current model:


- vary one of p, q, P, Q from current model by ± 1
- p, q both vary from current model by ± 1 .
- P, Q both vary from current model by ± 1 .
- Include/exclude c from current model

Model with lowest AIC becomes current model.

Repeat Step 2 until no lower AIC can be found.


ARIMA modelling

Forecasts from ARIMA(0,0,0)(1,0,0)[12] with non-zero mean



ARIMA modelling

Forecasts from ETS(M,Ad,M)



ARIMA vs ETS

- Myth that ARIMA models more general than exponential smoothing.

ARIMA vs ETS

- Myth that ARIMA models more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.

ARIMA vs ETS

- Myth that ARIMA models more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.

ARIMA vs ETS

- Myth that ARIMA models more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.
- Many ARIMA models which have no exponential smoothing counterparts.

ARIMA vs ETS

- Myth that ARIMA models more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.
- Many ARIMA models which have no exponential smoothing counterparts.
- ETS models all non-stationary. Models with seasonality or non-damped trend (or both) have two unit roots; all other models—that is, non-seasonal models with either no trend or damped trend—have one unit root.

Outline

- 1 Time series objects
- 2 Basic time series functionality
- 3 The forecast package
- 4 Exponential smoothing
- 5 ARIMA modelling
- 6 More from the forecast package**
- 7 Time series packages on CRAN

Other forecasting functions

- croston()** implements Croston's (1972) method for intermittent demand forecasting.
- theta()** provides forecasts from the Theta method.
- splinef()** gives cubic-spline forecasts, based on fitting a cubic spline to the historical data and extrapolating it linearly.
- meanf()** returns forecasts based on the historical mean.
- rwf()** gives "naïve" forecasts equal to the most recent observation assuming a random walk model.




Other plotting functions

tsdisplay() provides a time plot along with an ACF and PACF.


seasonplot() produces a seasonal plot.

tsdisplay

```
> tsdisplay(beer)
```



seasonplot



Outline

- 1 Time series objects
- 2 Basic time series functionality
- 3 The forecast package
- 4 Exponential smoothing
- 5 ARIMA modelling
- 6 More from the forecast package
- 7 Time series packages on CRAN**

Basic facilities

stats Contains substantial time series capabilities including the **ts** class for regularly spaced time series. Also ARIMA modelling, structural models, time series plots, acf and pacf graphs, classical decomposition and STL decomposition.

Forecasting and univariate modelling

- forecast** Lots of univariate time series methods including automatic ARIMA modelling, exponential smoothing via state space models, and the forecast class for consistent handling of time series forecasts. Part of the **forecasting** bundle.
- tseries** GARCH models and unit root tests.
- FitAR** Subset AR model fitting
- partsm** Periodic autoregressive time series models
- pear** Periodic autoregressive time series models

Forecasting and univariate modelling

Itsa Methods for linear time series analysis

dlm Bayesian analysis of Dynamic Linear Models.

timsac Time series analysis and control

fArma ARMA Modelling

fGarch ARCH/GARCH modelling

BootPR Bias-corrected forecasting and bootstrap prediction intervals for autoregressive time series

gsarima Generalized SARIMA time series simulation

bayesGARCH Bayesian Estimation of the GARCH(1,1) Model with t innovations

Resampling and simulation

boot Bootstrapping, including the block bootstrap with several variants.

meboot Maximum Entropy Bootstrap for Time Series

Decomposition and filtering

robfilter Robust time series filters

mFilter Miscellaneous time series filters useful for smoothing and extracting trend and cyclical components.

ArDec Autoregressive decomposition

wmtsa Wavelet methods for time series analysis based on Percival and Walden (2000)

wavelets Computing wavelet filters, wavelet transforms and multiresolution analyses

signalextraction Real-time signal extraction (direct filter approach)

bspec Bayesian inference on the discrete power spectrum of time series

Unit roots and cointegration

- tseries** Unit root tests and methods for computational finance.
- urca** Unit root and cointegration tests
- uroot** Unit root tests including methods for seasonal time series

Nonlinear time series analysis

nlts R functions for (non)linear time series analysis

tseriesChaos Nonlinear time series analysis

RTisean Algorithms for time series analysis from nonlinear dynamical systems theory.

tsDyn Time series analysis based on dynamical systems theory

BAYSTAR Bayesian analysis of threshold autoregressive models

fNonlinear Nonlinear and Chaotic Time Series Modelling

bentcableAR Bent-Cable autoregression

Dynamic regression models

- dynlm** Dynamic linear models and time series regression
- dyn** Time series regression
- tpr** Regression models with time-varying coefficients.

Multivariate time series models

mAr Multivariate AutoRegressive analysis

vars VAR and VEC models

MSBVAR Markov-Switching Bayesian Vector Autoregression Models

tsfa Time series factor analysis

dse Dynamic system equations including multivariate ARMA and state space models.

brainwaver Wavelet analysis of multivariate time series

Functional data

far Modelling Functional AutoRegressive processes

Continuous time data

- cts** Continuous time autoregressive models
- sde** Simulation and inference for stochastic differential equations.

Irregular time series

zoo Infrastructure for both regularly and irregularly spaced time series.

its Another implementation of irregular time series.

fCalendar Chronological and Calendarical Objects

fSeries Financial Time Series Objects

xts Provides for uniform handling of R's different time-based data classes

Time series data

fma Data from Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*. Part of the **forecasting** bundle.

expsmooth Data from Hyndman, Koehler, Ord and Snyder (2008) *Forecasting with exponential smoothing*. Part of the **forecasting** bundle.

Mcomp Data from the M-competition and M3-competition. Part of the **forecasting** bundle.

FinTS R companion to Tsay (2005) *Analysis of financial time series* containing data sets, functions and script files required to work some of the examples.

TSA R functions and datasets from Cryer and Chan (2008) *Time series analysis with applications in R*

TSdbi Common interface to time series databases

fame Interface for FAME time series databases

fEcofin Ecofin - Economic and Financial Data Sets

Miscellaneous

- hydrosanity** Graphical user interface for exploring hydrological time series
- pastecs** Regulation, decomposition and analysis of space-time series.
- RSEIS** Seismic time series analysis tools
- paleoTS** Modeling evolution in paleontological time-series
- GeneTS** Microarray Time Series and Network Analysis
- fractal** Fractal Time Series Modeling and Analysis