



# ETC3550: Applied forecasting for business and economics

Ch2. Time series graphics

[OTexts.org/fpp2/](https://OTexts.org/fpp2/)

# Outline

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Lag plots and autocorrelation
- 6 White noise

# ts objects and ts function

A time series is stored in a ts object in R:

- a list of numbers
- information about times those numbers were recorded.

## Example

Year	Observation
2012	123
2013	39
2014	78
2015	52
2016	110

```
y <- ts(c(123,39,78,52,110), start=2012)
```

## ts objects and ts function

For observations that are more frequent than once per year, add a frequency argument.

E.g., monthly data stored as a numerical vector z:

```
y <- ts(z, frequency=12, start=c(2003, 1))
```

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
--------------	-----------	---------------

Annual		
--------	--	--

Quarterly		
-----------	--	--

Monthly		
---------	--	--

Daily		
-------	--	--

Weekly		
--------	--	--

Hourly		
--------	--	--

Half-hourly		
-------------	--	--

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	
Quarterly		
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly		
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		



# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	
Weekly		
Hourly		
Half-hourly		

## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly		
Hourly		
Half-hourly		

## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	
Hourly		
Half-hourly		

## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly		
Half-hourly		

## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	
Half-hourly		



## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly		

## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly	48 or 336 or 17,532	

## ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly	48 or 336 or 17,532	1

# Australian GDP

```
ausgdp <- ts(x, frequency=4, start=c(1971,3))
```

- Class: "ts"


- Print and plotting methods available.

```
ausgdp
```

##	Qtr1	Qtr2	Qtr3	Qtr4
## 1971			4612	4651
## 1972	4645	4615	4645	4722
## 1973	4780	4830	4887	4933
## 1974	4921	4875	4867	4905
## 1975	4938	4934	4942	4979
## 1976	5028	5079	5112	5127
## 1977	5130	5101	5072	5069

# Australian GDP

```
autoplot(ausgdp)
```



# Residential electricity sales

```
elecsales
```

```
## Time Series:
```

```
## Start = 1989
```

```
## End = 2008
```

```
## Frequency = 1
```

```
## [1] 2354.34 2379.71 2318.52 2468.99 2386.09
```

```
## [6] 2569.47 2575.72 2762.72 2844.50 3000.70
```

```
## [11] 3108.10 3357.50 3075.70 3180.60 3221.60
```

```
## [16] 3176.20 3430.60 3527.48 3637.89 3655.00
```

# Class package

```
> library(fpp2)
```

# Class package

```
> library(fpp2)
```

This loads:

- some data for use in examples and exercises



# Class package

```
> library(fpp2)
```

This loads:


- some data for use in examples and exercises
- **forecast** package (for forecasting functions)
- **ggplot2** package (for graphics functions)
- **fma** package (for lots of time series data)
- **expsmooth** package (for more time series data)

# Outline

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Lag plots and autocorrelation
- 6 White noise


# Time plots

```
autoplot(melsyd[, "Economy.Class"])
```



# Time plots

```
autoplot(a10) + ylab("$ million") + xlab("Year") +  
  ggtitle("Antidiabetic drug sales")
```




## Your turn

- Create plots of the following time series: `dole`, `bricksq`, `lynx`, `goog`
- Use `help()` to find out about the data in each series.
- For the last plot, modify the axis labels and title.


# Are time plots best?

```
autoplot(elecdaily[, "Temperature"]) +  
  xlab("Week") + ylab("Max temperature")
```




# Are time plots best?

```
qplot(time(elecdaily), elecdaily[, "Temperature"]) +  
  xlab("Week") + ylab("Max temperature")
```

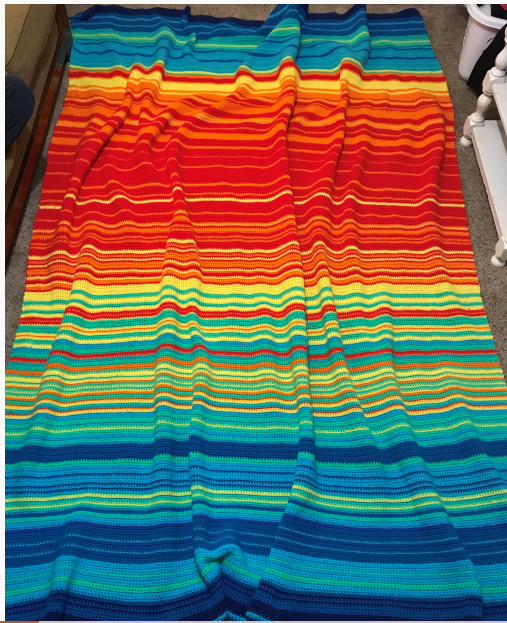


# Are time plots best?





# Are time plots best?




# Outline

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Lag plots and autocorrelation
- 6 White noise

# Seasonal plots

```
ggseasonplot(a10, year.labels=TRUE, year.labels.left=TRUE) +  
  ylab("$ million") +  
  ggtitle("Seasonal plot: antidiabetic drug sales")
```




# Seasonal plots

- Data plotted against the individual “seasons” in which the data were observed. (In this case a “season” is a month.)
- Something like a time plot except that the data from each season are overlapped.
- Enables the underlying seasonal pattern to be seen more clearly, and also allows any substantial departures from the seasonal pattern to be easily identified.
- In R: `ggseasonplot()`

# Seasonal polar plots


```
ggseasonplot(a10, polar=TRUE) + ylab("$ million")
```

Seasonal plot: a10



# Seasonal subseries plots

```
ggsubseriesplot(a10) + ylab("$ million") +  
  ggtitle("Subseries plot: antidiabetic drug sales")
```




## Seasonal subseries plots

- Data for each season collected together in time plot as separate time series.
- Enables the underlying seasonal pattern to be seen clearly, and changes in seasonality over time to be visualized.
- In R: `ggsubseriesplot()`

# Quarterly Australian Beer Production


```
beer <- window(ausbeer, start=1992)  
autoplot(beer)
```






# Quarterly Australian Beer Production

```
ggseasonplot(beer, year.labels=TRUE)
```



# Quarterly Australian Beer Production

```
ggsubseriesplot(beer)
```



## Your turn

The `arrivals` data set comprises quarterly international arrivals (in thousands) to Australia from Japan, New Zealand, UK and the US.

- Use `autoplot()` and `ggseasonplot()` to compare the differences between the arrivals from these four countries.
- Can you identify any unusual observations?

# Outline

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Lag plots and autocorrelation
- 6 White noise

# Time series patterns

**Trend** pattern exists when there is a long-term increase or decrease in the data.

**Seasonal** pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).

**Cyclic** pattern exists when data exhibit rises and falls that are *not of fixed period* (duration usually of at least 2 years).


# Time series components

## Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern


# Time series patterns

```
autoplot(window(elec, start=1980)) +  
  ggtitle("Australian electricity production") +  
  xlab("Year") + ylab("GWh")
```



# Time series patterns


```
autoplot(bricksq) +  
  ggtitle("Australian clay brick production") +  
  xlab("Year") + ylab("million units")
```






# Time series patterns

```
autoplot(hsales) +  
  ggtitle("Sales of new one-family houses, USA") +  
  xlab("Year") + ylab("Total sales")
```




# Time series patterns

```
autoplot(ustreas) +  
  ggtitle("US Treasury Bill Contracts") +  
  xlab("Day") + ylab("price")
```



# Time series patterns

```
autoplot(lynx) +  
  ggtitle("Annual Canadian Lynx Trappings") +  
  xlab("Year") + ylab("Number trapped")
```



# Seasonal or cyclic?

## Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

# Seasonal or cyclic?

## Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

The timing of peaks and troughs is predictable with seasonal data, but unpredictable in the long term with cyclic data.

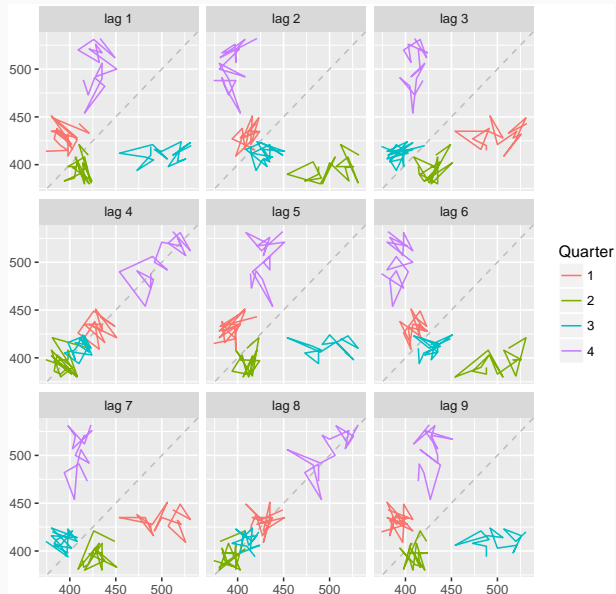
# Outline

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Lag plots and autocorrelation
- 6 White noise

## Example: Beer production

```
beer <- window(ausbeer, start=1992)  
gglagplot(beer)
```

# Example: Beer production





# Lagged scatterplots

- Each graph shows  $y_t$  plotted against  $y_{t-k}$  for different values of  $k$ .
- The autocorrelations are the correlations associated with these scatterplots.

# Autocorrelation

**Covariance** and **correlation**: measure extent of **linear relationship** between two variables ( $y$  and  $X$ ).

# Autocorrelation

**Covariance** and **correlation**: measure extent of **linear relationship** between two variables ( $y$  and  $X$ ).

**Autocovariance** and **autocorrelation**: measure linear relationship between **lagged values** of a time series  $y$ .

# Autocorrelation

**Covariance** and **correlation**: measure extent of **linear relationship** between two variables ( $y$  and  $X$ ).

**Autocovariance** and **autocorrelation**: measure linear relationship between **lagged values** of a time series  $y$ .

We measure the relationship between:

- $y_t$  and  $y_{t-1}$
- $y_t$  and  $y_{t-2}$
- $y_t$  and  $y_{t-3}$
- etc.

# Autocorrelation

We denote the sample autocovariance at lag  $k$  by  $c_k$  and the sample autocorrelation at lag  $k$  by  $r_k$ . Then define

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and  $r_k = c_k / c_0$

# Autocorrelation

We denote the sample autocovariance at lag  $k$  by  $c_k$  and the sample autocorrelation at lag  $k$  by  $r_k$ . Then define

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and  $r_k = c_k / c_0$


- $r_1$  indicates how successive values of  $y$  relate to each other
- $r_2$  indicates how  $y$  values two periods apart relate to each other
- $r_k$  is *almost* the same as the sample correlation between  $y_t$  and  $y_{t-k}$ .

# Autocorrelation

Results for first 9 lags for beer data:

$r_1$	$r_2$	$r_3$	$r_4$	$r_5$	$r_6$	$r_7$	$r_8$	$r_9$
-0.102	-0.657	-0.060	0.869	-0.089	-0.635	-0.054	0.832	-0.108

`ggAcf(beer)`




# Autocorrelation

- $r_4$  higher than for the other lags. This is due to **the seasonal pattern in the data**: the peaks tend to be **4 quarters** apart and the troughs tend to be **2 quarters** apart.
- $r_2$  is more negative than for the other lags because troughs tend to be 2 quarters behind peaks.
- Together, the autocorrelations at lags 1, 2, ..., make up the *autocorrelation* or ACF.
- The plot is known as a **correlogram**



```
ggAcf (beer)
```




# Trend and seasonality in ACF plots

- When data have a trend, the autocorrelations for small lags tend to be large and positive.
- When data are seasonal, the autocorrelations will be larger at the seasonal lags (i.e., at multiples of the seasonal frequency)
- When data are trended and seasonal, you see a combination of these effects.


# Aus monthly electricity production

```
elec2 <- window(elec, start=1980)  
autoplot(elec2)
```



# Aus monthly electricity production

```
ggAcf(elec2, lag.max=48)
```



# Aus monthly electricity production


Time plot shows clear trend and seasonality.

The same features are reflected in the ACF.

- The slowly decaying ACF indicates trend.
- The ACF peaks at lags 12, 24, 36, ..., indicate seasonality of length 12.


# Google stock price

```
autoplot(goog)
```



# Google stock price

```
ggAcf(goog, lag.max=100)
```



# Your turn

We have introduced the following graphics functions:

- `gglagplot`
- `ggAcf`


Explore the following time series using these functions. Can you spot any seasonality, cyclicity and trend? What do you learn about the series?

- `hsales`
- `usdeaths`
- `bricksq`
- `sunspotarea`
- `gasoline`




# Which is which?


1. Daily temperature of cow




2. Monthly accidental deaths




3. Monthly air passengers




4. Annual mink trappings




A




B



C



D




# Outline

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Lag plots and autocorrelation
- 6 White noise

# Example: White noise


```
wn <- ts(rnorm(36))  
autoplot(wn)
```



# Example: White noise

---

$r_1$	-0.04
$r_2$	-0.01
$r_3$	-0.15
$r_4$	-0.11
$r_5$	-0.06
$r_6$	-0.01
$r_7$	0.01
$r_8$	0.06
$r_9$	0.22
$r_{10}$	-0.12



Sample autocorrelations for white noise series.

We expect each autocorrelation to be close to zero.

# Sampling distribution of autocorrelations

Sampling distribution of  $r_k$  for white noise data is asymptotically  $N(0, 1/T)$ .

# Sampling distribution of autocorrelations

Sampling distribution of  $r_k$  for white noise data is asymptotically  $N(0, 1/T)$ .

- 95% of all  $r_k$  for white noise must lie within  $\pm 1.96/\sqrt{T}$ .
- If this is not the case, the series is probably not WN.
- Common to plot lines at  $\pm 1.96/\sqrt{T}$  when plotting ACF. These are the **critical values**.


# Autocorrelation

Example:

$T = 36$  and so critical values at

$$\pm 1.96 / \sqrt{36} = \pm 0.327.$$

All autocorrelation coefficients lie within these limits, confirming that the data are white noise. (More precisely, the data cannot be distinguished from white noise.)



# Example: Pigs slaughtered


```
pigs2 <- window(pigs, start=1990)
autoplot(pigs2) +
  xlab("Year") + ylab("thousands") +
  ggtitle("Number of pigs slaughtered in Victoria")
```





# Example: Pigs slaughtered

```
ggAcf(pigs2)
```



## Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)

## Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows some significant autocorrelation at lags 1, 2, and 3.
- $r_{12}$  relatively large although not significant. This may indicate some slight seasonality.

## Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows some significant autocorrelation at lags 1, 2, and 3.
- $r_{12}$  relatively large although not significant. This may indicate some slight seasonality.

These show the series is **not a white noise series**.

## Your turn

You can compute the daily changes in the Google stock price using

```
dgoog <- diff(goog)
```

Does dgoog look like white noise?