

# Forecasting: principles and practice

**Rob J Hyndman**

1.1 Time series graphics

# Outline

- 1 Time series in R**
- 2 Time plots
- 3 Lab session 1
- 4 Seasonal plots
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation
- 7 White noise
- 8 Lab session 2

# ts objects and ts function

A time series is stored in a ts object in R:

- a list of numbers
- information about times those numbers were recorded.

## Example

Year	Observation
2012	123
2013	39
2014	78
2015	52
2016	110

```
y <- ts(c(123,39,78,52,110), start=2012)
```

# ts objects and ts function

For observations that are more frequent than once per year, add a frequency argument.

E.g., monthly data stored as a numerical vector z:

```
y <- ts(z, frequency=12, start=c(2003, 1))
```

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start	example
--------------	-----------	-------	---------

Annual

Quarterly

Monthly

Daily

Weekly

Hourly

Half-hourly

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	
Quarterly		
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly		
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		



# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly		
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily		
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly		
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly		
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	
Half-hourly		



# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly		

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly	48 or 336 or 17,532	

# ts objects and ts function

```
ts(data, frequency, start)
```

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	1 or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly	48 or 336 or 17,532	1

# Australian GDP

```
ausgdp <- ts(scan("gdp.dat"),frequency=4,  
             start=1971+2/4)
```


- Class: "ts"
- Print and plotting methods available.

```
ausgdp
```

```
##           Qtr1 Qtr2 Qtr3 Qtr4  
## 1971           4612 4651  
## 1972 4645 4615 4645 4722  
## 1973 4780 4830 4887 4933  
## 1974 4921 4875 4867 4905  
## 1975 4938 4934 4942 4979
```

# Australian GDP

```
autoplot(ausgdp)
```



# Residential electricity sales

```
elecsales
```

```
## Time Series:
```

```
## Start = 1989
```

```
## End = 2008
```

```
## Frequency = 1
```

```
## [1] 2354.34 2379.71 2318.52 2468.99 2386.09
```

```
## [9] 2844.50 3000.70 3108.10 3357.50 3075.70
```


```
## [17] 3430.60 3527.48 3637.89 3655.00
```

# Outline

- 1 Time series in R
- 2 Time plots**
- 3 Lab session 1
- 4 Seasonal plots
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation
- 7 White noise
- 8 Lab session 2

# Time plots


```
autoplot(melsyd[, "Economy.Class"])
```





# Time plots

```
autoplot(a10) + ylab("$ million") + xlab("Year") +  
  ggtitle("Antidiabetic drug sales")
```



# Outline

- 1 Time series in R
- 2 Time plots
- 3 Lab session 1**
- 4 Seasonal plots
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation
- 7 White noise
- 8 Lab session 2

# Lab Session 1


# Outline

- 1 Time series in R
- 2 Time plots
- 3 Lab session 1
- 4 Seasonal plots**
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation
- 7 White noise
- 8 Lab session 2

# Seasonal plots

```
ggseasonplot(a10, ylab="$ million",  
  year.labels=TRUE, year.labels.left=TRUE) +  
  ggtitle("Seasonal plot: antidiabetic drug sales")
```

Seasonal plot: antidiabetic drug sales




# Seasonal plots

- Data plotted against the individual “seasons” in which the data were observed. (In this case a “season” is a month.)
- Something like a time plot except that the data from each season are overlapped.
- Enables the underlying seasonal pattern to be seen more clearly, and also allows any substantial departures from the seasonal pattern to be easily identified.
- In R: `ggseasonplot`

# Seasonal polar plots

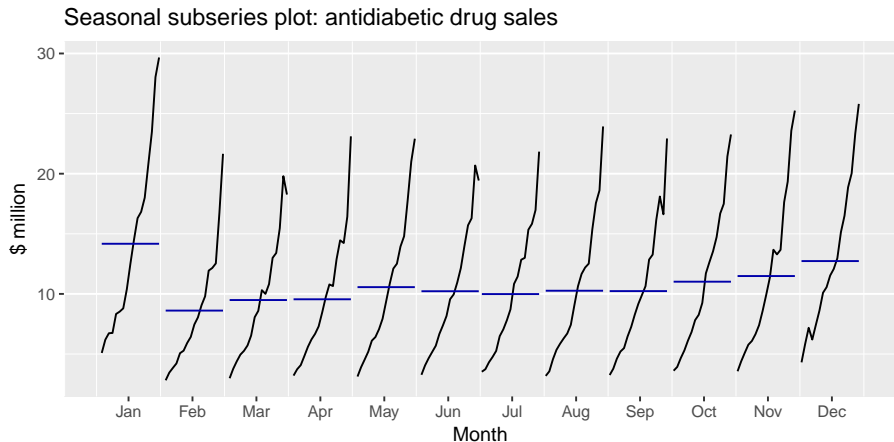
```
ggseasonplot(a10, polar=TRUE) + ylab("$ million")
```

Seasonal plot: a10



# Seasonal subseries plots

```
ggmonthplot(a10) + ylab("$ million") +  
  ggtitle("Seasonal subseries plot: antidiabetic drug sales")
```






# Seasonal subseries plots

- Data for each season collected together in time plot as separate time series.
- Enables the underlying seasonal pattern to be seen clearly, and changes in seasonality over time to be visualized.
- In R: `ggmonthplot`

# Quarterly Australian Beer Production


```
beer <- window(ausbeer, start=1992)  
autoplot(beer)
```



# Quarterly Australian Beer Production


```
ggseasonplot(beer, year.labels=TRUE)
```

Seasonal plot: beer



# Quarterly Australian Beer Production

```
ggsubseriesplot(beer)
```



# Outline

- 1 Time series in R
- 2 Time plots
- 3 Lab session 1
- 4 Seasonal plots
- 5 Seasonal or cyclic?**
- 6 Lag plots and autocorrelation
- 7 White noise
- 8 Lab session 2

# Time series patterns

**Trend** pattern exists when there is a long-term increase or decrease in the data.

**Seasonal** pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).

**Cyclic** pattern exists when data exhibit rises and falls that are *not of fixed period* (duration usually of at least 2 years).

# Time series components


## Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

# Time series patterns

```
autoplot(window(elec, start=1980)) +  
  ggtitle("Australian electricity production")  
  xlab("Year") + ylab("GWh")
```


Australian electricity production





# Time series patterns


```
autoplot(bricksq) +  
  ggtitle("Australian clay brick production") +  
  xlab("Year") + ylab("million units")
```



# Time series patterns

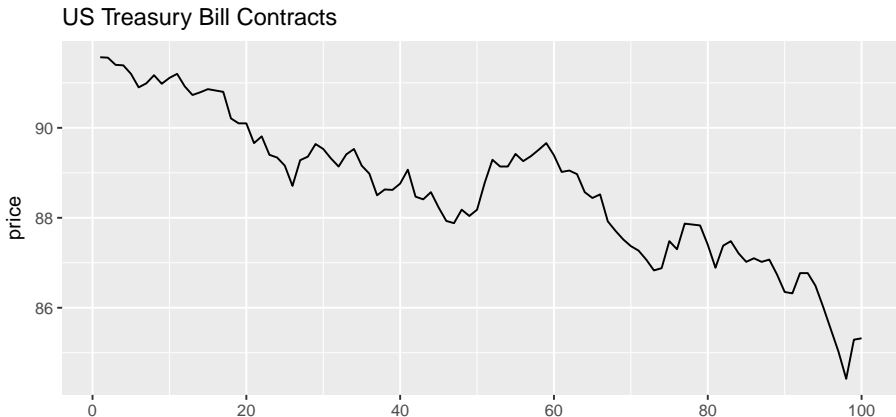
```
autoplot(hsales) +  
  ggtitle("Sales of new one-family houses, USA")  
  xlab("Year") + ylab("Total sales")
```

Sales of new one-family houses, USA



# Time series patterns


```
autoplot(ustreas) +  
  ggtitle("US Treasury Bill Contracts") +  
  xlab("Day") + ylab("price")
```



# Time series patterns

```
autoplot(lynx) +  
  ggtitle("Annual Canadian Lynx Trappings") +  
  xlab("Year") + ylab("Number trapped")
```

Annual Canadian Lynx Trappings



# Seasonal or cyclic?

## Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

The timing of peaks and troughs is predictable with seasonal data, but unpredictable in the long term with cyclic data.

# Seasonal or cyclic?

## Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

The timing of peaks and troughs is predictable with seasonal data, but unpredictable in the long term with cyclic data.

# Outline

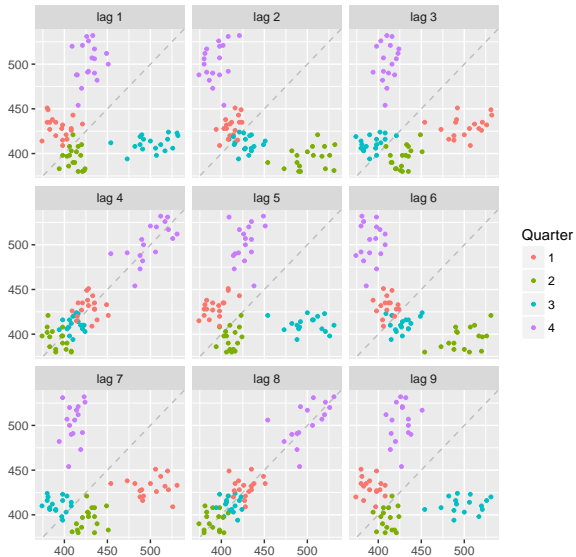
- 1 Time series in R
- 2 Time plots
- 3 Lab session 1
- 4 Seasonal plots
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation**
- 7 White noise
- 8 Lab session 2

# Example: Beer production

```
beer <- window(ausbeer, start=1992)
gglagplot(beer, lags=9, do.lines=FALSE,
           continuous=FALSE)
```



# Example: Beer production



# Lagged scatterplots

- Each graph shows  $y_t$  plotted against  $y_{t-k}$  for different values of  $k$ .
- The autocorrelations are the correlations associated with these scatterplots.

# Autocorrelation

**Covariance** and **correlation**: measure extent of **linear relationship** between two variables ( $y$  and  $X$ ).

**Autocovariance** and **autocorrelation**: measure linear relationship between **lagged values** of a time series  $y$ .

We measure the relationship between:

- $y_t$  and  $y_{t-1}$
- $y_t$  and  $y_{t-2}$
- $y_t$  and  $y_{t-3}$
- etc.

# Autocorrelation

**Covariance** and **correlation**: measure extent of **linear relationship** between two variables ( $y$  and  $X$ ).

**Autocovariance** and **autocorrelation**: measure linear relationship between **lagged values** of a time series  $y$ .

We measure the relationship between:

- $y_t$  and  $y_{t-1}$
- $y_t$  and  $y_{t-2}$
- $y_t$  and  $y_{t-3}$
- etc.

# Autocorrelation

**Covariance** and **correlation**: measure extent of **linear relationship** between two variables ( $y$  and  $X$ ).

**Autocovariance** and **autocorrelation**: measure linear relationship between **lagged values** of a time series  $y$ .

We measure the relationship between:

- $y_t$  and  $y_{t-1}$
- $y_t$  and  $y_{t-2}$
- $y_t$  and  $y_{t-3}$
- etc.

# Autocorrelation

We denote the sample autocovariance at lag  $k$  by  $c_k$  and the sample autocorrelation at lag  $k$  by  $r_k$ . Then define

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and  $r_k = c_k / c_0$

- $r_1$  indicates how successive values of  $y$  relate to each other
- $r_2$  indicates how  $y$  values two periods apart relate to each other
- $r_k$  is *almost* the same as the sample correlation between  $y_t$  and  $y_{t-k}$ .

# Autocorrelation

We denote the sample autocovariance at lag  $k$  by  $c_k$  and the sample autocorrelation at lag  $k$  by  $r_k$ . Then define

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and  $r_k = c_k / c_0$


- $r_1$  indicates how successive values of  $y$  relate to each other
- $r_2$  indicates how  $y$  values two periods apart relate to each other
- $r_k$  is *almost* the same as the sample correlation between  $y_t$  and  $y_{t-k}$ .

# Autocorrelation

Results for first 9 lags for beer data:/footnotesize

$r_1$	$r_2$	$r_3$	$r_4$	$r_5$	$r_6$	$r_7$	$r_8$	$r_9$
-0.102	-0.657	-0.060	0.869	-0.089	-0.635	-0.054	0.832	-0.102

```
ggAcf(beer)
```





# Autocorrelation

- $r_4$  higher than for the other lags. This is due to **the seasonal pattern in the data**: the peaks tend to be **4 quarters** apart and the troughs tend to be **2 quarters** apart.
- $r_2$  is more negative than for the other lags because troughs tend to be 2 quarters behind peaks.
- Together, the autocorrelations at lags 1, 2, ..., make up the *autocorrelation* or ACF.
- The plot is known as a **correlogram**


# Recognizing seasonality in a time series

If there is seasonality, the ACF at the seasonal lag (e.g., 12 for monthly data) will be **large and positive**.

- For seasonal monthly data, a large ACF value will be seen at lag 12 and possibly also at lags 24, 36, ...
- For seasonal quarterly data, a large ACF value will be seen at lag 4 and possibly also at lags 8, 12, ...


# Aus monthly electricity production

```
elec2 <- window(elec, start=1980)  
autoplot(elec2)
```



# Aus monthly electricity production

```
ggAcf(elec2, lag.max=48)
```



# Aus monthly electricity production


Time plot shows clear trend and seasonality.

The same features are reflected in the ACF.


- The slowly decaying ACF indicates trend.
- The ACF peaks at lags 12, 24, 36, ..., indicate seasonality of length 12.

# Which is which?


1. Daily temperature of cow




2. Monthly accidental deaths




3. Monthly air passengers




4. Annual mink trappings




A




B



C



D




# Outline

- 1 Time series in R
- 2 Time plots
- 3 Lab session 1
- 4 Seasonal plots
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation
- 7 White noise**
- 8 Lab session 2

# Example: White noise


```
wn <- ts(rnorm(36))  
autoplot(wn)
```





# Example: White noise


```
wn <- ts(rnorm(36))  
autoplot(wn)
```



# Example: White noise

---

$r_1$	-0.09
$r_2$	-0.02
$r_3$	0.14
$r_4$	-0.00
$r_5$	-0.06
$r_6$	-0.08
$r_7$	0.14
$r_8$	-0.10
$r_9$	-0.21
$r_{10}$	-0.04



Sample autocorrelations for white noise series.

For uncorrelated data, we would expect each autocorrelation to be close to zero.

# Sampling distribution of autocorrelations

Sampling distribution of  $r_k$  for white noise data is asymptotically  $N(0, 1/T)$ .

- 95% of all  $r_k$  for white noise must lie within  $\pm 1.96/\sqrt{T}$ .
- If this is not the case, the series is probably not WN.
- Common to plot lines at  $\pm 1.96/\sqrt{T}$  when plotting ACF. These are the *critical values*.

# Sampling distribution of autocorrelations

Sampling distribution of  $r_k$  for white noise data is asymptotically  $N(0, 1/T)$ .

- 95% of all  $r_k$  for white noise must lie within  $\pm 1.96/\sqrt{T}$ .
- If this is not the case, the series is probably not WN.
- Common to plot lines at  $\pm 1.96/\sqrt{T}$  when plotting ACF. These are the *critical values*.


# Autocorrelation

## Example:

$T = 36$  and so critical values at

$$\pm 1.96 / \sqrt{36} = \pm 0.327.$$


All autocorrelation coefficients lie within these limits, confirming that the data are white noise. (More precisely, the data cannot be distinguished from white noise.)



# Example: Pigs slaughtered


```
pigs2 <- window(pigs, start=1990)
autoplot(pigs2) +
  xlab("Year") + ylab("thousands") +
  ggtitle("Number of pigs slaughtered in Victoria")
```

Number of pigs slaughtered in Victoria



# Example: Pigs slaughtered

```
ggAcf(pigs2)
```



# Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows some significant autocorrelation at lags 1, 2, and 3.
- $r_{12}$  relatively large although not significant. This may indicate some slight seasonality.

These show the series is **not a white noise series**.



# Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows some significant autocorrelation at lags 1, 2, and 3.
- $r_{12}$  relatively large although not significant. This may indicate some slight seasonality.

These show the series is **not a white noise series**.

# Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)




- Difficult to detect pattern in time plot.
- ACF shows some significant autocorrelation at lags 1, 2, and 3.
- $r_{12}$  relatively large although not significant. This may indicate some slight seasonality.

These show the series is **not a white noise series**.

# Combination graph

```
ggtsdisplay(pigs2, plot.type='scatter')
```

pigs2



# Outline

- 1 Time series in R
- 2 Time plots
- 3 Lab session 1
- 4 Seasonal plots
- 5 Seasonal or cyclic?
- 6 Lag plots and autocorrelation
- 7 White noise
- 8 Lab session 2**

# Lab Session 2